

RL-glue and you

Why RL-glue?

- In RL we have agent's interacting with an environment
 - e.g., a bandit algorithm interacting with a 10-armed bandit problem
 - e.g., a Sarsa agent interacting with a grid world
- The agent and environment are both **programs** that must *exchange* data at each time step
- RL-glue manages this *interaction*, deciding many things: like who goes first? what information is stored?, etc

What would you do?

- Given the task of running an experiment with an agent and environment?
- You could do the *single-file approach*: write one big program that contains the agent code, the environment code, and some function to call each of them in turn
 - What are the limitations of this approach?
- You could make your own agent, environment software, architecture.

What is RL-Glue

- RL-glue is a *communication protocol*
 - *a system of rules that allow two or more entities to transmit information*
 - *...defining the rules syntax, semantics and synchronization of communication*

A standard interface

- The implementation of the RL-glue protocol is realized as a standard software **interface** that the agent and environment **must implement**
- An agent that is compatible with RL-glue is required to implement a set of standard functions:
 - *agent_init, agent_start, agent_step, agent_cleanup, agent_message*
 - a python program implements these functions is compatible with RL-glue. We call it an ***agent program***

Standard interface...

- An environment that is compatible with RL-glue is required to implement a set of standard functions:
 - *env_init, env_start, env_step, env_cleanup, env_message*
- a python program implements these functions is compatible with RL-glue. We call it an **environment program**

simple_agent.py

- We released some **example** code to show you how to write agents, and environments compatible with RL-glue
- simple_agent.py is an example of one such agent. It's a simple agent that does not learn, but instead takes random actions
- but it does provide an **example** of an agent that is compatible with RL-glue: it implements the *agent interface* (**all the required agent functions!**)

simple_env.py

- simple_env.py is an **example** of an environment. It ignores the agent's action and returns a random reward on each time step
- But it does provide an **example** of an environment that is compatible with RL-glue: it implements the *environment interface* (**all the required functions!**)

Your task

- Make your own agent and environment programs:
 - Maybe you might call them: `bandit_agent.py`, and `bandit_env.py`
- Start by copying over the code in `simple_agent.py` and `simple_env.py`
 - that way you know you have implemented the required functions
- Modify the functions inside `bandit_agent.py` so that it implements the bandit algorithm described in the assignment
 - similarly for the `bandit_env.py`

Let's look at simple agent

```
#!/usr/bin/env python
```

```
Comments
```

```
imports ...
```

```
globals ...
```

```
def agent_init():
```

```
    # this is where you initialize agent data structures
```

```
def agent_start(this_observation):
```

```
    # first time step
```

```
    # this is where your agent selects an action and returns it. No learning happens here!
```

```
    return action
```

```
def agent_step(reward, this_observation):
```

```
    # the agent selects a new action
```

```
    # and does a learning update with the Q-values, based on the reward and action from previous time step
```

```
    return action
```

```
def agent_end(reward):
```

```
    # final learning update at end of episode
```

```
    # never gets called in a bandit problem
```

```
    return
```

```
def agent_cleanup():
```

```
    # clean up
```

```
    return
```

```
def agent_message(inMessage):
```

```
    # might be useful to get information from the agent
```

```
    return "some string"
```

Study simple_env.py

- It tells you what functions your bandit_env.py program should implement
- Can you figure out from simple_env.py and the README file what each function does?

Experiment programs

- Now we have an agent and environment program that are compatible with RL-glue!
 - How do we run an experiment?
- We must write an experiment program!
- The experiment program **calls functions** inside rl_glue.py
- These rl_glue functions provide different ways to run an experiment

**Let's look at simple experiment
(an Example experiment program)**

```
#!/usr/bin/env python
```

```
Comments
```

```
imports ...
```

```
from rl_glue import * # Required for RL-Glue
```

```
RLGlue("simple_env", "simple_agent")
```

```
if __name__ == "__main__":
```

```
    maxSteps = 1000
```

```
    numRuns = 30
```

```
    # create some data structures to store performance
```

```
    for k in range(numRuns):
```

```
        RL_init()
```

```
        RL_start()
```

```
        for i in range(maxSteps):
```

```
            (reward, obs, action, terminal) = RL_step()
```

```
            # store some of this information
```

```
        RL_cleanup()
```

```
        # update statistics about agent performance, so we can average over runs
```

```
    # Save data to a file
```

What can you modify?

- Anything inside your agent, environment, and experiment programs:
 - `bandit_agent.py`, `bandit_env.py`, `bandit_exp.py`
- Your agent **must** implement all of the required agent functions
- And your environment **must** implement all the required environment functions
- And your experiment program **should** call functions in `rl_glue.py`

What can you **not** modify?

- Don't touch rl_glue.py
- There will be a few lines of code that are special
 - for example, the experiment program must contain the following lines of code:
 - `from rl_glue import * # Required for RL-Glue`
`RLGlue("simple_env", "simple_agent")`
 - These cases are well documented in the example code!

Exercise: find the glue!!

A simple bandit algorithm

Initialize, for $a = 1$ to k :

$$Q(a) \leftarrow 0$$

$$N(a) \leftarrow 0$$

Repeat forever:

$$A \leftarrow \begin{cases} \arg \max_a Q(a) & \text{with probability } 1 - \varepsilon \quad (\text{breaking ties randomly}) \\ \text{a random action} & \text{with probability } \varepsilon \end{cases}$$

$$R \leftarrow \text{bandit}(A)$$

$$N(A) \leftarrow N(A) + 1$$

$$Q(A) \leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)]$$

Agent program

Environment program

Experiment program