


# What we learned last time

- Value-function approximation by stochastic gradient descent enables RL to be applied to arbitrarily large state spaces
- Most algorithms just carry over the Targets from the tabular case
- With bootstrapping (TD), we don't get true gradient descent methods
  - but the linear, on-policy case is still guaranteed convergent
  - and learning is *faster* with  $n$ -step methods ( $n > 1$ ), as before
- For continuous state spaces, coarse/tile coding is a good strategy

Chapter 10:

# On-policy Control with Approximation

Value function approximation (VFA) replaces the table with a general parameterized form



# On-policy Control with Approximation

- (Semi-)gradient methods carry over to control in the usual way
  - Mountain Car example
- $n$ -step methods carry over too, with the usual tradeoffs
- A new *average-reward* setting,  
with *differential* value functions and *differential* algorithms
  - Queuing example (tabular)
- The discounting setting is deprecated

# (Semi-)gradient methods carry over to control in the usual on-policy GPI way

- Always learn the action-value function of the current policy
- Always act near-greedily wrt the current action-value estimates
- The learning rule is the same as in Chapter 9:

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha \left[ U_t - \hat{q}(S_t, A_t, \boldsymbol{\theta}_t) \right] \nabla \hat{q}(S_t, A_t, \boldsymbol{\theta}_t)$$

update target, e.g.,  $U_t = G_t$  (MC)

$U_t = R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, \boldsymbol{\theta}_t)$  (Sarsa)

(Expected Sarsa)  $U_t = R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) \hat{q}(S_{t+1}, a, \boldsymbol{\theta}_t)$        $U_t = \sum_{s', r} p(s', r|S_t, A_t) \left[ r + \gamma \sum_{a'} \pi(a'|s') \hat{q}(s', a', \boldsymbol{\theta}_t) \right]$  (DP)

# (Semi-)gradient methods carry over to control

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha \left[ U_t - \hat{q}(S_t, A_t, \boldsymbol{\theta}_t) \right] \nabla \hat{q}(S_t, A_t, \boldsymbol{\theta}_t)$$

## Episodic Semi-gradient Sarsa for Estimating $\hat{q} \approx q_*$

Input: a differentiable function  $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^n \rightarrow \mathbb{R}$

Initialize value-function weights  $\boldsymbol{\theta} \in \mathbb{R}^n$  arbitrarily (e.g.,  $\boldsymbol{\theta} = \mathbf{0}$ )

Repeat (for each episode):

$S, A \leftarrow$  initial state and action of episode (e.g.,  $\varepsilon$ -greedy)

Repeat (for each step of episode):

Take action  $A$ , observe  $R, S'$

If  $S'$  is terminal:

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \alpha [R - \hat{q}(S, A, \boldsymbol{\theta})] \nabla \hat{q}(S, A, \boldsymbol{\theta})$$

Go to next episode


Choose  $A'$  as a function of  $\hat{q}(S', \cdot, \boldsymbol{\theta})$  (e.g.,  $\varepsilon$ -greedy)

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \alpha [R + \gamma \hat{q}(S', A', \boldsymbol{\theta}) - \hat{q}(S, A, \boldsymbol{\theta})] \nabla \hat{q}(S, A, \boldsymbol{\theta})$$

$$S \leftarrow S'$$

$$A \leftarrow A'$$

# Example: The Mountain-Car problem



Minimum-Time-to-Goal Problem

SITUATIONS:

car's position and velocity

ACTIONS:


three thrusts: forward, reverse, none

REWARDS:

always  $-1$  until car reaches the goal


Episodic, No Discounting,  $\gamma=1$

# Values learned while solving Mountain-Car with tile coding function approximation





# Learning curves for semi-gradient Sarsa with tile coding

Mountain Car  
Steps per episode  
log scale  
averaged over 100 runs



# $n$ -step semi-gradient Sarsa is better for $n > 1$

$$\theta_{t+n} \doteq \theta_{t+n-1} + \alpha \left[ G_t^{(n)} - \hat{q}(S_t, A_t, \theta_{t+n-1}) \right] \nabla \hat{q}(S_t, A_t, \theta_{t+n-1}), \quad 0 \leq t < T$$



# On-policy Control with Approximation

- (Semi-)gradient methods carry over to control in the usual way
  - Mountain Car example
- $n$ -step methods carry over too, with the usual tradeoffs
  - A new average-reward setting, with differential value functions and differential algorithms
    - Queuing example (tabular)
  - The discounting setting is deprecated

# On-policy Control with Approximation

- (Semi-)gradient methods carry over to control in the usual way
  - Mountain Car example
- $n$ -step methods carry over too, with the usual tradeoffs
- A new average-reward setting,  
with differential value functions and differential algorithms
  - Queuing example (tabular)
- The discounting setting is deprecated

# A new goal for continuing tasks: Maximizing average reward per time step

$$\begin{aligned}\text{Maximize } \eta(\pi) &\doteq \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \mathbb{E}[R_t \mid A_{0:t-1} \sim \pi] && \text{assuming that these limits exist} \\ &= \lim_{t \rightarrow \infty} \mathbb{E}[R_t \mid A_{0:t-1} \sim \pi], && \text{is known as the } \textit{ergodicity} \text{ property} \\ &= \sum_s d_\pi(s) \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)r\end{aligned}$$

$d_\pi : \mathcal{S} \rightarrow [0, 1]$  is the steady-state distribution under  $\pi$ ,  
also known as the on-policy distribution:

$$d_\pi(s) \doteq \lim_{t \rightarrow \infty} \Pr\{S_t = s \mid A_{0:t-1} \sim \pi\}$$

$\eta(\pi)$  is the *average* amount of *reward* received per time step

# In the average reward setting, everything is new



- Returns:  $G_t \doteq R_{t+1} - \eta(\pi) + R_{t+2} - \eta(\pi) + R_{t+3} - \eta(\pi) + \dots$
- Bellman Eqs:  $v_\pi(s) = \sum_a \pi(a|s) \sum_{r,s'} p(s', r|s, a) [r - \eta(\pi) + v_\pi(s')],$  prediction  
 $q_\pi(s, a) = \sum_{r,s'} p(s', r|s, a) [r - \eta(\pi) + \sum_{a'} \pi(a'|s') q_\pi(s', a')],$   
 $v_*(s) = \max_a \sum_{r,s'} p(s', r|s, a) [r - \eta(\pi) + v_*(s')],$  and control  
 $q_*(s, a) = \sum_{r,s'} p(s', r|s, a) [r - \eta(\pi) + \max_{a'} q_*(s', a')]$
- Update targets:

$$U_t \doteq R_{t+1} - \bar{R}_t + \hat{q}(S_{t+1}, A_{t+1}, \theta) \quad \text{or} \quad U_t \doteq R_{t+1} - \bar{R}_t + \hat{v}(S_{t+1}, \theta)$$

 estimate of  $\eta(\pi)$

## Differential semi-gradient Sarsa for estimating $\hat{q} \approx q_*$

Input: a differentiable function  $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^n \rightarrow \mathbb{R}$

Parameters: step sizes  $\alpha, \beta > 0$

Initialize value-function weights  $\boldsymbol{\theta} \in \mathbb{R}^n$  arbitrarily (e.g.,  $\boldsymbol{\theta} = \mathbf{0}$ )

Initialize average reward estimate  $\bar{R}$  arbitrarily (e.g.,  $\bar{R} = 0$ )

Initialize state  $S$ , and action  $A$

Repeat (for each step):

    Take action  $A$ , observe  $R, S'$

    Choose  $A'$  as a function of  $\hat{q}(S', \cdot, \boldsymbol{\theta})$  (e.g.,  $\varepsilon$ -greedy)

$\delta \leftarrow R - \bar{R} + \hat{q}(S', A', \boldsymbol{\theta}) - \hat{q}(S, A, \boldsymbol{\theta})$

$\bar{R} \leftarrow \bar{R} + \beta \delta$


$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \alpha \delta \nabla \hat{q}(S, A, \boldsymbol{\theta})$

$S \leftarrow S'$

$A \leftarrow A'$

# Example: The access-control queuing problem solved by tabular differential Sarsa

- Customers wait in line to be served by one of  $k=10$  servers
- Customers pay rewards of 1, 2, 4, or 8 (depending on their priority) for being served
- On each step, the customer at the front of the queue is accepted (served), or rejected
- The queue never empties; new customers have random priorities
- Busy servers become free with probability  $p=0.06$  on each step



$$t = 2,000,000, \alpha = \beta = .01, \epsilon = .1, \bar{R}_t \approx 2.31$$

# Discounting is futile in continuing control settings with function approximation

- The problem statement is broken! The goal is broken!
- We can no longer give a useful ordering on policies
  - we can only order a few policies, those that dominate others in all states
  - It would be OK if we could say what states we care about, but in the control case we can't
  - Suppose we cared about states according to how often they occur? Surprisingly, discounting then becomes irrelevant!



## The Futility of Discounting in Continuing Problems

Perhaps discounting can be saved by choosing an objective that sums discounted values over the distribution with which states occur under the policy:

$$\begin{aligned} J(\pi) &= \sum_s d_\pi(s) v_\pi^\gamma(s) && \text{(where } v_\pi^\gamma \text{ is the discounted value function)} \\ &= \sum_s d_\pi(s) \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r|s, a) [r + \gamma v_\pi^\gamma(s')] && \text{(Bellman Eq.)} \\ &= \eta(\pi) + \sum_s d_\pi(s) \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r|s, a) \gamma v_\pi^\gamma(s') && \text{(from (10.5))} \\ &= \eta(\pi) + \gamma \sum_{s'} v_\pi^\gamma(s') \sum_s d_\pi(s) \sum_a \pi(a|s) p(s'|s, a) && \text{(from (3.8))} \\ &= \eta(\pi) + \gamma \sum_{s'} v_\pi^\gamma(s') d_\pi(s') && \text{(from (10.6))} \\ &= \eta(\pi) + \gamma J(\pi) \\ &= \eta(\pi) + \gamma \eta(\pi) + \gamma^2 J(\pi) \\ &= \eta(\pi) + \gamma \eta(\pi) + \gamma^2 \eta(\pi) + \gamma^3 \eta(\pi) + \dots \\ &= \frac{1}{1 - \gamma} \eta(\pi). \end{aligned}$$

The proposed discounted objective orders policies identically to the undiscounted (average reward) objective. We have failed to save discounting!

# Conclusions

- Control is straightforward in the on-policy, episodic, linear case
- For the continuing case, we need the average-reward setting
  - which is a lot like just replacing  $R_t$  with  $R_t - \eta(\pi)$  everywhere
  - where  $\eta(\pi)$  is the average reward per step, or its estimate
- We should probably never use discounting as a control objective
- Formal results (bounds) exist for the linear, on-policy case
  - we get chattering near a good solution, not convergence