

Deterministic Tree Search

aka Deterministic Tree-based Planning
aka "Search"

finding a path from start state to a goal state

Search Algorithm Properties


- **Complete?** Guaranteed to find a solution if one exists?
- **Optimal?** Guaranteed to find the least cost path?
- **Time complexity?**
- **Space complexity?**

Variables:

n	Number of states in the problem
b	The average branching factor B (the average number of successors)
s	Depth of the shallowest solution
m	Max depth of the search tree

Depth-First Search (DFS)

- With cycle checking, DFS is complete.




Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$

- When is DFS optimal?

Breadth-First Search (BFS)

Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$
BFS		Y	N*	$O(b^{s+1})$	$O(b^s)$



- When is BFS optimal?


Comparisons

- When will BFS outperform DFS?
- When will DFS outperform BFS?

Iterative Deepening


Iterative deepening uses DFS as a subroutine:

1. Do a DFS which only searches for paths of length 1 or less.
 2. If “1” failed, do a DFS which only searches paths of length 2 or less.
 3. If “2” failed, do a DFS which only searches paths of length 3 or less.
-and so on.



Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$
BFS		Y	N*	$O(b^{s+1})$	$O(b^s)$
ID		Y	N*	$O(b^{s+1})$	$O(bs)$

Costs on Actions





Notice that BFS finds the shortest path in terms of number of transitions. It does not find the least-cost path.

We will now cover an algorithm which does find the least-cost path.

Uniform Cost Search


Expand cheapest node first:

Fringe is a priority queue



Uniform Cost Search



Algorithm		Complete	Optimal	Time	Space
DFS	w/ Path Checking	Y	N	$O(b^{m+1})$	$O(bm)$
BFS		Y	N	$O(b^{s+1})$	$O(b^s)$
UCS		Y*	Y	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$



** UCS can fail if actions can get arbitrarily cheap*

Uniform Cost Issues

- Remember: explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
 - Explores options in every “direction”
 - No information about goal location
 - “blind” search



Recap: Search

- **Search problem:**
 - States (configurations of the world)
 - Deterministic transitions: a function from states to lists of (next state, cost) pairs; drawn as a graph
 - Start state and goal test
- **Search tree:**
 - Nodes: represent plans for reaching states
 - Plans have costs (sum of action costs)
- **Search Algorithm:**
 - Systematically builds a search tree
 - Chooses an ordering of the fringe (unexplored nodes)




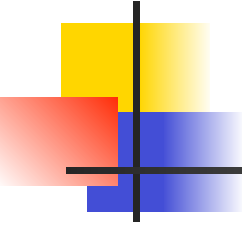
Game-tree search

- Approximate
 - Evaluation functions
 - Anytime algorithms
- Adversarial
 - Minimax search
 - Alpha-Beta pruning

based on work by R Greiner, D Lin, Jean-Claude
Latombe, N Nilsson

Partial Game Search Tree for Tic-Tac-Toe





But... in general the search tree is too big to make it possible to reach the terminal states!

Examples:

- Checkers: $\sim 10^{40}$ nodes
- Chess: $\sim 10^{120}$ nodes




Evaluation Function of a State

- $e(s) = +\infty$ if s is a win for MAX
- $e(s) = -\infty$ if s is a win for MIN
- $e(s)$ = a measure of how “favorable”
is s for MAX
 - > 0 if s is considered favorable to MAX
 - < 0 otherwise




Example: Tic-Tac-Toe


$e(s)$ = number of rows, columns, and diagonals open for MAX
- number of rows, columns, and diagonals open for MIN



$$8-8 = 0$$



$$6-4 = 2$$



$$3-3 = 0$$



Evaluation Function for chess

- $e(s)$ = weighted sum of feature

$$e(s) \doteq \boldsymbol{\theta}^\top \boldsymbol{\phi} \doteq \theta_1 \phi_1 + \theta_2 \phi_2 + \cdots + \theta_n \phi_n$$

- Features


- # of white pawns – # of black pawns
- # of white bishops – # of black bishops
- # of white rooks – # of black rooks
- ...

- Weights

- 1 for pawns, 3 for bishops, 5 for rooks

Example

Tic-Tac-Toe with
horizon = 2






Minimax

- Achieves “Perfect” play for deterministic, perfect-information games
- Idea: choose move leading to position with highest minimax value
 - best achievable payoff against best play

Example: a 2-ply game





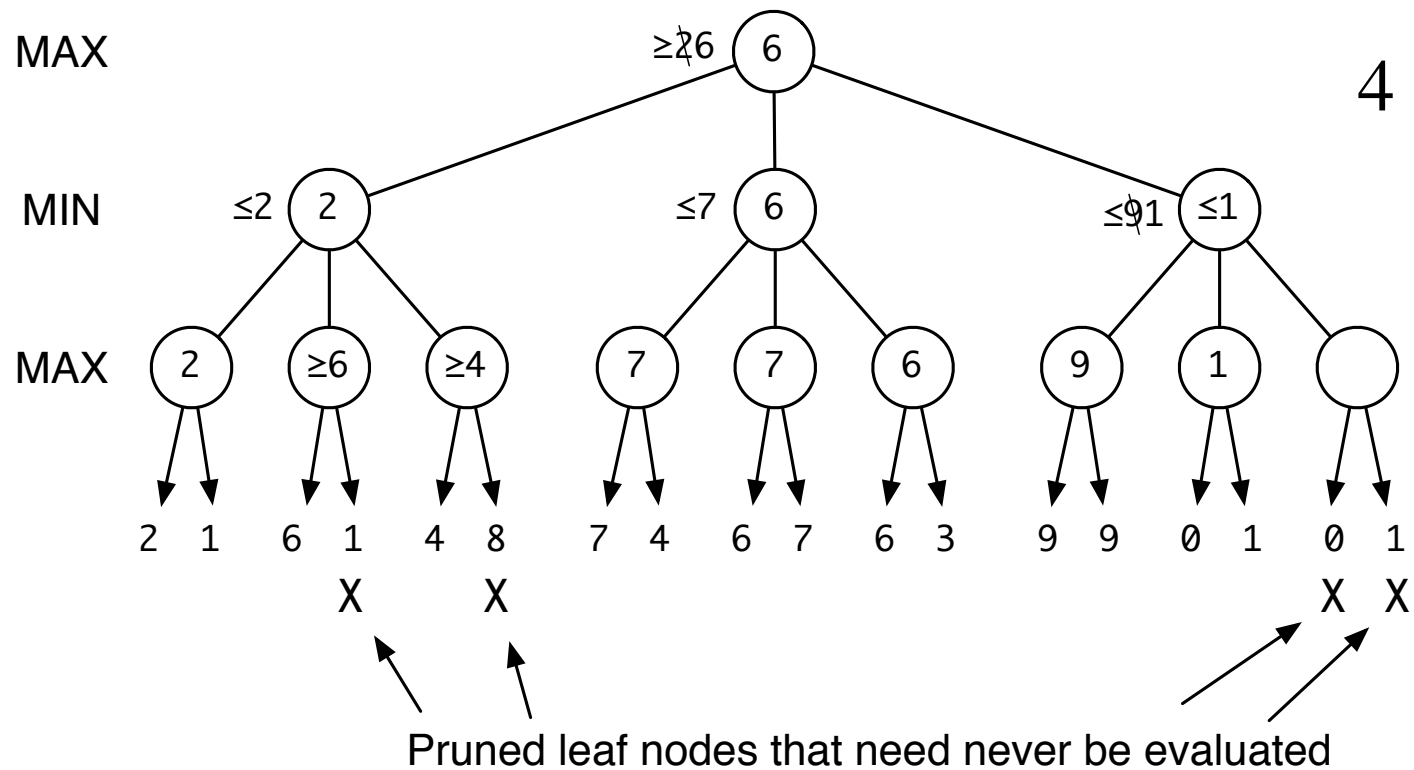
Minimax (back of the envelope)

- Does minimax work in practice?


In chess: we can do about $b^d = 10^6$
thus, if $b=35$, then $d=4$
but 4-ply lookahead is a hopeless chess
player!

- 4-ply = human novice
8-ply = typical PC, human master
12-ply = Deep Blue, Kasparov

α-β Pruning



Alpha-beta pruning greatly thins the tree





Properties of α - β Search

- Pruning does not affect final result
- Good move ordering improves effectiveness of pruning
- With “perfect ordering”:
 - time complexity = $O(b^{d/2})$
 - doubles depth of search
 - can easily reach depth 8
 - play good chess!



Summary

- There are lots of different kinds of search
 - With different optimizations and guarantees
- All involve planning – using your knowledge of the world's dynamics to anticipate the consequences of your action, and then picking the best
- All search involves computing how good it is to be in each state
 - And benefits from a good initial estimate