

Heuristic Single-Agent Search


Robert Holte



A Heuristic Function Estimates Distance to Goal



Heuristics Speed up Search



10,461,394,944,000 states

heuristic search examines 36,000

Example Heuristic Functions

- $h(n)$ estimates cost of cheapest path from node n to goal node
- Example: 8-puzzle

5		8
4	2	1
7	3	6

n

1	2	3
4	5	6
7	8	

goal

$h_1(n)$ = number of misplaced tiles
= 6

Example Heuristic Functions

- $h(n)$ estimates cost of cheapest path from node n to goal node
- Example: 8-puzzle

5		8
4	2	1
7	3	6

n

1	2	3
4	5	6
7	8	

goal


$h_1(n)$ = number of misplaced tiles
= 6

$h_2(n)$ = sum of the distances of
every tile to its goal position
= 3 + 1 + 3 + 0 + 2 + 1 + 0 + 3
= 13

How to Use a Heuristic Function to Speed Up Search ?



Notation



- $g(n)$ = distance from start to node n along our current path (not necessarily optimal)
- $h(n)$ = estimated distance from n to goal
- $f(n) = g(n) + h(n)$ = estimated distance from start to goal via n (using our current path to n)

More Notation

- $g^*(n)$ = true distance from start to node n
- $h^*(n)$ = true distance from n to goal
- $f^*(n)$ = true distance from start to goal via n
= $g^*(n) + h^*(n)$

How to Use a Heuristic Function to Speed Up Search ?



(1) Pruning

- If we have a bound B on total solution cost (e.g. iterative deepening) we can prune n from the search space if $f(n) > B$.
- But this might prune all the optimal paths. One condition that makes pruning safe is ...
- $g^*(n) + h(n) \leq f^*(n)$, which is true if $h(n)$...
- never overestimates, $h(n) \leq h^*(n)$.
- Such a heuristic is called “admissible”.

Iterative Deepening Heuristic Search (IDA*)

- Using this pruning idea, when you generate a node n you only explore beneath it (with a recursive call) if $f(n) \leq B$, the current cost bound (not depth bound).
- How do you update the cost bound from one iteration to the next ?

(2) Node Ordering

Use the heuristic to decide which node to expand next. Two main options.

1. Choose the node with minimum $h(n)$.
Uniform-cost search using this selection rule is called “Pure Heuristic Search”
2. Choose the node with minimum $f(n)$.
Uniform-cost search using this selection rule is called A^* .

Terminology

- Expanding a state means generating its successors (children).
- A state is open if it has been generated but not expanded.
- A state is closed if it has been expanded.
- Open list = data structure holding all currently open states.
- Closed list = data structure indicating which states are closed.

Dijkstra's Algorithm*

- Put (start,0) in OPEN
- Repeat
 - If OPEN is empty, exit with failure
 - Select (n,g(n)) on OPEN with minimum g(n)
 - If n is the goal, return path from start to goal
 - Move (n,g(n)) from OPEN to CLOSED
 - For each successor x of n:
 - If (x,g(x)) is on OPEN and $g(x) > g(n) + \text{cost}(n,x)$ then replace (x,g(x)) on OPEN with (x,g(n)+cost(n,x))
 - If (x,_) is on neither OPEN nor CLOSED then add (x,g(n)+cost(n,x)) to OPEN

* Uniform-cost Search

A* Algorithm

- Put (start, $f(\text{start})$) in OPEN
- Repeat
 - If OPEN is empty, exit with failure
 - Select $(n, f(n))$ on OPEN with minimum $f(n)$
 - If n is the goal, return path from start to goal
 - Move $(n, f(n))$ from OPEN to CLOSED
 - For each successor x of n :
 - If $(x, f(x))$ is on OPEN and $f(x) > g(n) + \text{cost}(n, x) + h(x)$ then replace $(x, f(x))$ on OPEN with $(x, g(n) + \text{cost}(n, x) + h(x))$
 - If $(x, _)$ is on neither OPEN nor CLOSED then add $(x, g(n) + \text{cost}(n, x) + h(x))$ to OPEN
 - If $(x, f(x))$ is on CLOSED and $f(x) > g(n) + \text{cost}(n, x) + h(x)$ then remove $(x, f(x))$ from CLOSED and add $(x, g(n) + \text{cost}(n, x) + h(x))$ to OPEN


Question

When these algorithms stop (first remove the goal from Open), under what conditions are we guaranteed that this path to the goal is optimal ?

Dijkstra: all edge weights are non-negative

A*: impose conditions on the heuristic

Condition on the Heuristic



Require $\langle N, g^*(N) + h(N) \rangle$ lower cost than $\langle G, P \rangle$


$$g^*(N) + h(N) < P$$

$$\Leftrightarrow h(N) < P - g^*(N)$$

$$\Leftarrow h(N) \leq h^*(N) \quad (\text{because } h^*(N) < P - g^*(N))$$

Admissible Heuristic \Rightarrow A^* stops with an optimal path to goal

A* must re-open closed nodes




OPEN: (G,140), (C,90)

CLOSED: (S,70), (B,40), (C,110), (D,110), ...(A,120)

How Bad Can This Be ?

- If N is the number of distinct nodes that A^* expands, the total number of node expansions could be as bad as $O(2^N)$.
- Some A^* variations reduce the worst-case to $O(N^2)$.
- Alternatively, if the heuristic has a special property, basic A^* never re-opens a node once it is closed, so is $O(N)$.

Consistent Heuristic



Require $\langle N, g^*(N) + h(N) \rangle$ lower cost than $\langle X, P + h(X) \rangle$

$$g^*(N) + h(N) < P + h(X)$$

$$\Leftrightarrow h(N) - h(X) < P - g^*(N)$$

$$\Leftarrow h(N) - h(X) \leq c(N, X)$$

(because $c(N, X) < P - g^*(N)$)

A heuristic is consistent if $h(N) \leq c(N, X) + h(X)$
for all N and **all** X.

Consistent \Rightarrow first path to X off Open is optimal for all X

The Need for Automatically Generated Heuristics





Summary

- There are lots of different kinds of search
 - With different optimizations and guarantees
- All involve planning – using your knowledge of the world's dynamics to anticipate the consequences of your action, and then picking the best
- All search involves computing how good it is to be in each state
 - And benefits from a good initial estimate